

Insufficient number of features:  
The Abalone dataset contains only 8 input features, which may not be sufficient to capture all the complexity in the data. If the input features do not provide enough information to accurately predict the target variable, the model may not perform well.

Changing the parameter  $n\_estimator=100$ ,  $lr=0.1$ , and max depth 3 for raw abalone dataset: Accuracy score: 0.2548

Upon using more optimum parameters for Gradient Boosting, the accuracy increases. This low accuracy may be due to the fact that the features are highly correlated.

### Gradient Boosting on Wine dataset

Accuracy score on Raw wine dataset: 59.92

The accuracy of Gradient boosting on the wine - raw dataset is more than that of Random forests and this may be due to the fact that the dataset has outliers and is not balanced. When the dataset contains imbalanced classes, Random Forests may produce biased predictions towards the majority class, as each tree is built independently and can be influenced by the class imbalance, while Gradient Boosting Classifier can adjust the weights of the samples to balance the classes.

### Gradient Boosting on Abalone - PCA dataset

Accuracy score: 0.1148

Classification Report:				
	precision	recall	f1-score	support
3	0.00	0.00	0.00	3
4	0.29	0.15	0.20	13
5	0.04	0.03	0.03	32
6	0.05	0.04	0.05	48
7	0.00	0.00	0.00	84
8	0.10	0.01	0.02	99
9	0.10	0.06	0.07	142
10	0.18	0.54	0.27	139
11	0.00	0.00	0.00	93
12	0.25	0.02	0.04	51
13	0.04	0.06	0.05	31
14	0.00	0.00	0.00	26
15	0.00	0.00	0.00	21
16	0.03	0.15	0.05	13
17	0.00	0.00	0.00	8
18	0.03	0.17	0.05	12
19	0.00	0.00	0.00	7
20	0.00	0.00	0.00	4
21	0.00	0.00	0.00	3
22	0.00	0.00	0.00	3
23	0.00	0.00	0.00	4
accuracy			0.11	836
macro avg	0.05	0.06	0.04	836
weighted avg	0.08	0.11	0.07	836

#### Note:

The accuracy on PCA dataset upon using Gradient Boosting is lesser than Random forests. Overall it can be seen that PCA hurts the performance of a tree boosting classifier as data has been lost while reducing the number of dimensions.

```
Confusion Matrix:
[[ 0  3  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  2  0  0  1  0  0  10  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  2  1  0  0  0  1  28  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  6  2  0  1  16  20  0  0  2  0  0  0  0  1  0  0  0  0]
 [ 0  0 11  9  0  2  26  28  0  1  2  0  0  0  0  5  0  0  0  0]
 [ 0  0  0  8  3  1  15  42  0  0 18  0  0  1  0 11  0  0  0  0]
 [ 0  0  6  6 14  3  8  79  0  1  8  0  0  2  0 15  0  0  0  0]
 [ 0  0  2  3 15  1  7  75  0  1 10  0  0 13  0 11  0  0  1  0]
 [ 0  0  1  3 18  1  4  39  0  0  0  0  0 19  0  6  0  0  2  0]
 [ 0  0  0  3  4  1  2  26  0  1  1  0  0  9  0  4  0  0  0  0]
 [ 0  0  0  0  2  0  3  16  0  0  2  0  0  6  0  2  0  0  0  0]
 [ 0  0  0  1  1  0  0 12  0  0  3  0  0  4  0  5  0  0  0  0]
 [ 0  0  0  1  0  0  0 16  0  0  0  0  0  2  0  2  0  0  0  0]
 [ 0  0  0  0  0  0  0  9  0  0  1  0  0  2  0  0  0  0  1  0]
 [ 0  0  0  0  0  0  0  4  0  0  1  0  0  3  0  0  0  0  0  0]
 [ 0  0  0  0  3  0  0  7  0  0  0  0  0  0  0  2  0  0  0  0]
 [ 0  0  0  0  0  0  0  4  0  0  0  0  0  3  0  0  0  0  0  0]
 [ 0  0  0  0  1  0  0  2  0  0  0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  1  0  0  0  0  0  0  0  0  0  1  0  1  0  0  0  0]
 [ 0  0  0  0  0  0  0  2  0  0  0  0  0  0  0  0  0  0  1  0]
 [ 0  0  0  0  1  0  0  2  0  0  1  0  0  0  0  0  0  0  0  0]]
Training Score: 0.10356180784196349
Testing Score: 0.11483253588516747
```

### Gradient Boosting on Wine - PCA dataset

Accuracy score on wine dataset using PCA preprocessing: 0.5400

Classification Report:				
	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.32	0.22	0.26	46
5	0.57	0.56	0.57	420
6	0.56	0.63	0.59	579
7	0.50	0.38	0.43	221
8	0.35	0.25	0.29	32
accuracy			0.54	1300
macro avg	0.38	0.34	0.36	1300
weighted avg	0.54	0.54	0.54	1300

#### Note:

The training score is 82% whereas the same classifier has a training score of approximately 70% on raw data without PCA reduction. So in this case, PCA helps in improving the accuracy but there is a considerable amount of overfitting.

```
Confusion Matrix:
[[ 0  0  2  0  0  0]
 [ 1 10 10 21  2  2]
 [ 2 10 237 158 10  3]
 [ 3  7 141 362 61  5]
 [ 1  3  24 103 85  5]
 [ 0  1  3  8 12  8]]
Training Score: 0.8289397729459304
Testing Score: 0.54
```

Accuracy score: 0.2105

Accuracy score: 0.5531

**Note:**

Training score is high as when the dataset has a small number of samples, Gradient boosting can overfit and since most features in the abalone dataset is highly correlated, dimensionality reduction has a positive effect on efficient computation.

But testing score is very low as there is considerable loss of data and Gradient boosting works better with more features.

The mean accuracy using Random Forests is 0.27 whereas for Gradient boosting, it is lower. This is possible if there are too many outliers/high correlation in the dataset, which is true for this case.

**Note:**

There is less overfitting in the training data after using LDA and Gradient boosting techniques. The test accuracy is also close but not very high. Compared to random forests, the accuracy is similar on Wine - LDA dataset.